AutoML – The rise of the machines

Nuno César de Sá[1],

[1] Institute of Environmental Sciences (CML)

AutoML – The rise of the machines

- Machine learning
  - What is it, and how does it work?
- AutoML
- “hands-on” example on RS data
  - Optional, naturally

CCTV footage of an AutoML system dealing with a GIS expert

Is this machine learning?

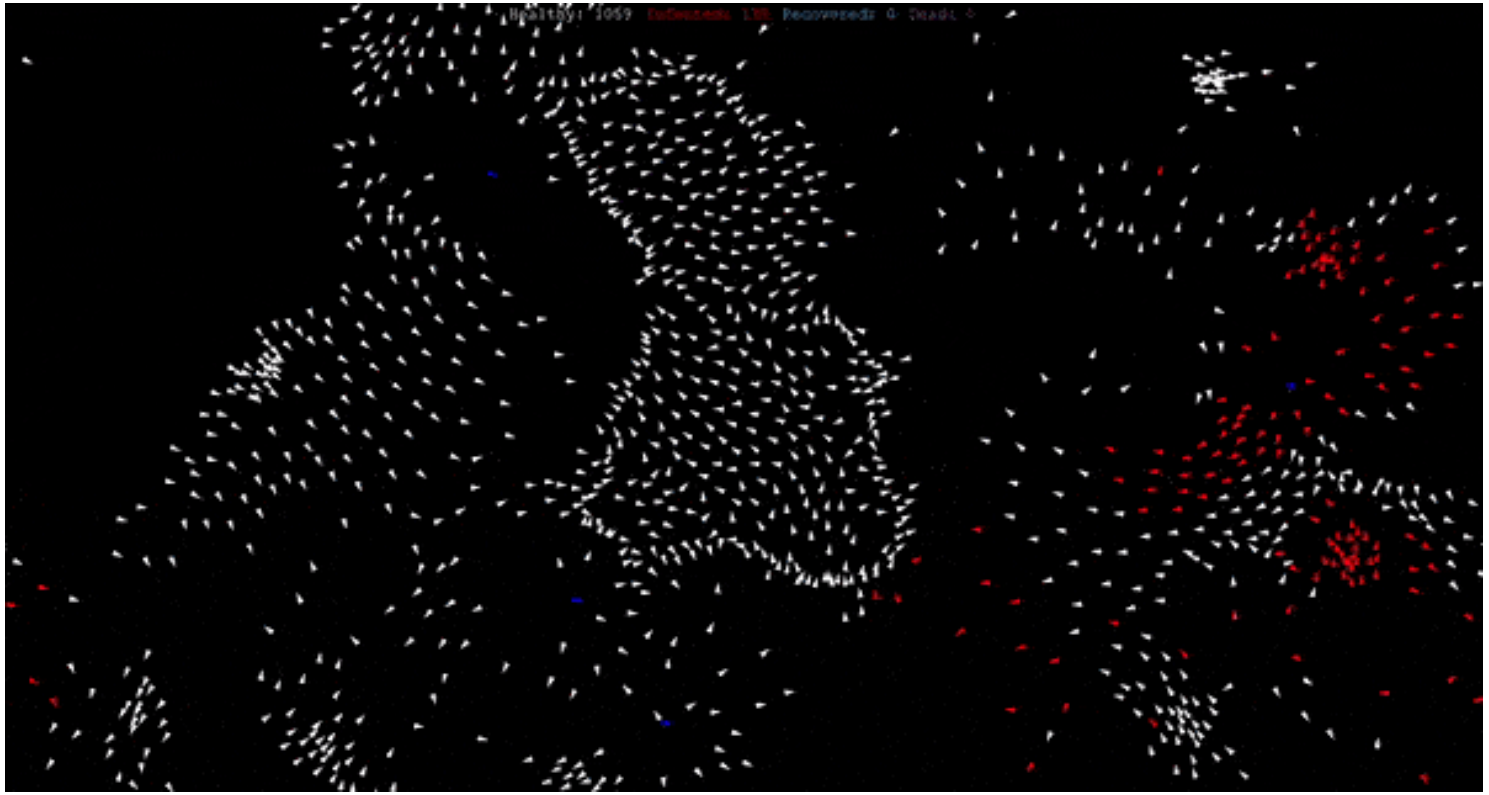

Flocking algorithm (Reynolds, 1986) – more info @ link

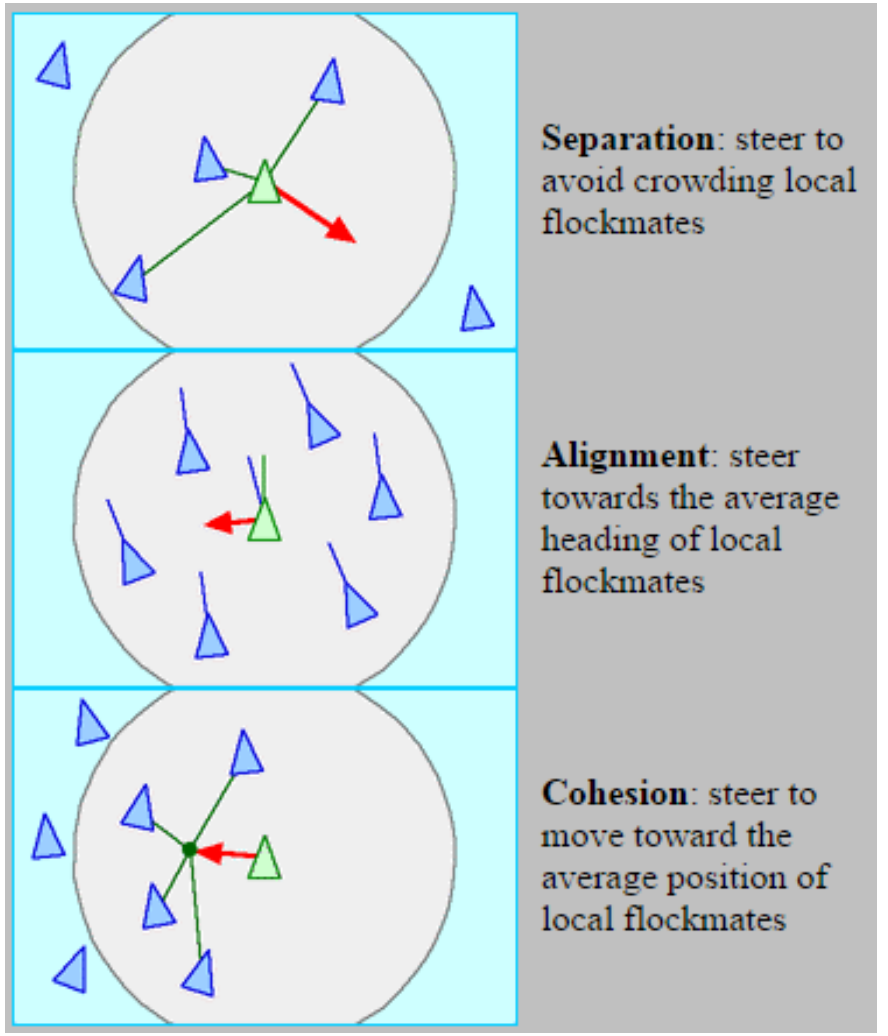

Is this machine learning?

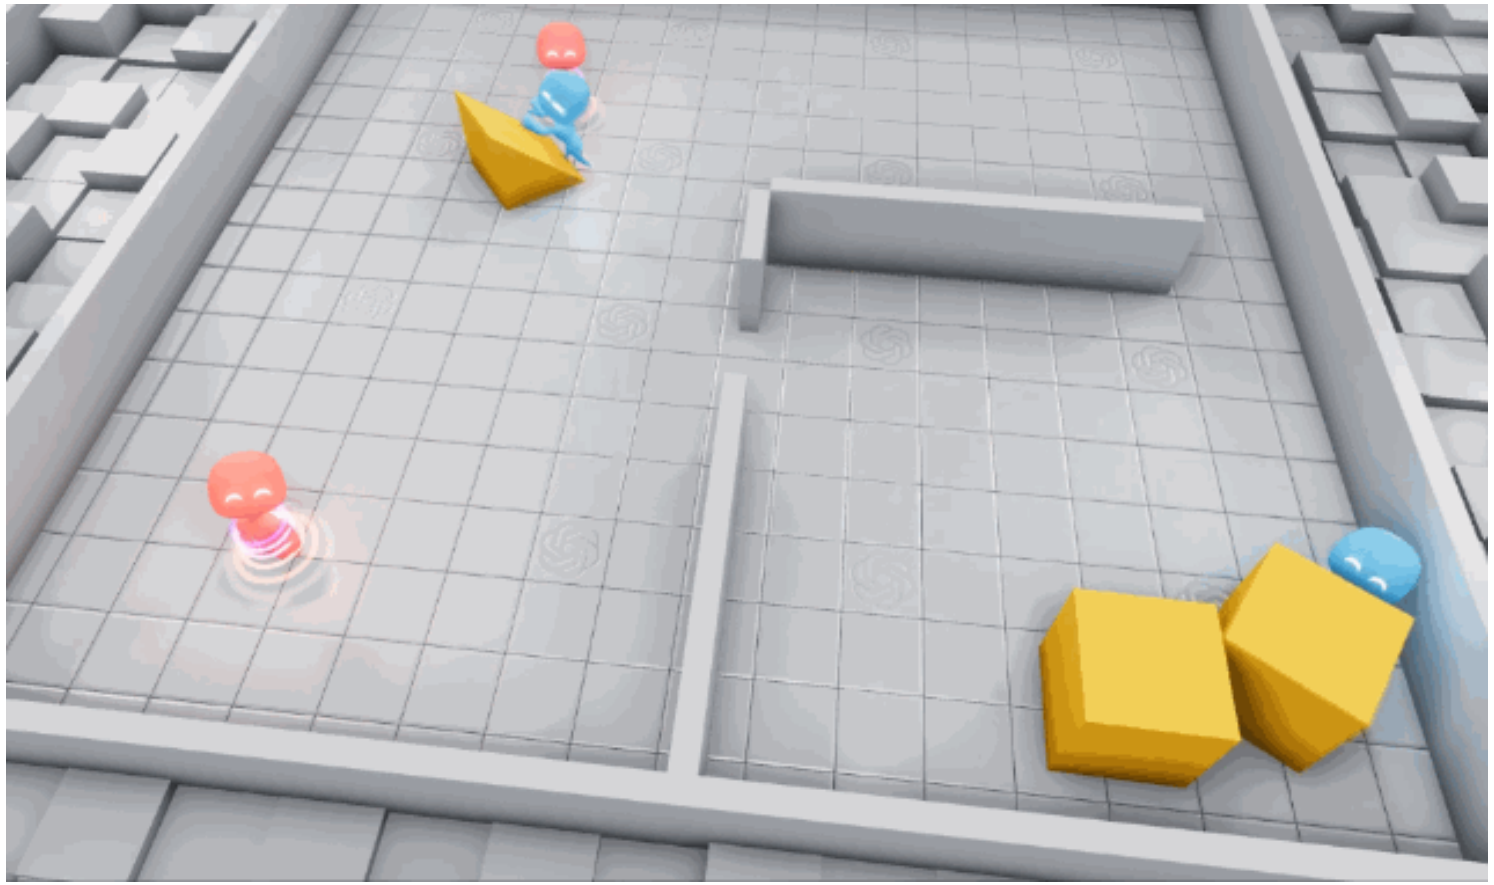

Emergent Tool Use From Multi-Agent Autocurricula, Bowen baker et al; 2 minutes papers

- Hide & Seek game:
  - Reds are seekers
  - Blues are hiding
- Agents:
  - Move blocks
  - Have a field of vision
  - Measure distance to everything else
- Survival strategies are learned

Machine learning – (mis)conceptions

- What is machine learning?
  - Arthur Samuel: "the field of study that gives computers the ability to learn without being explicitely programmed"

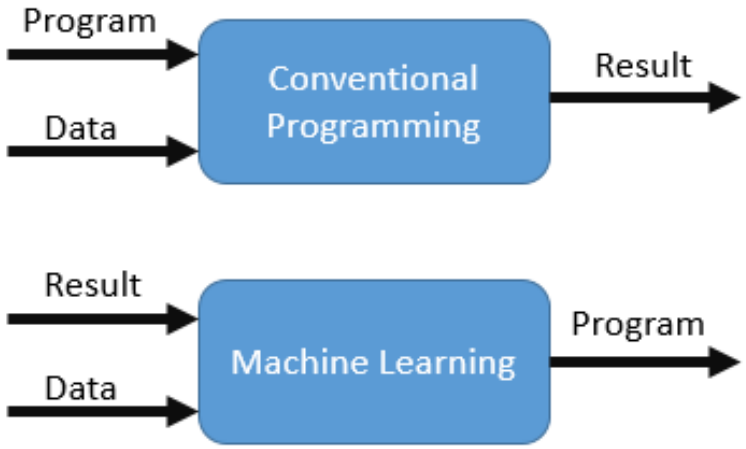

- It's not "new"
  - First concepts of AI – *Boomer* technology
  - "Black-Box" -> mostly advanced statistics + linear models

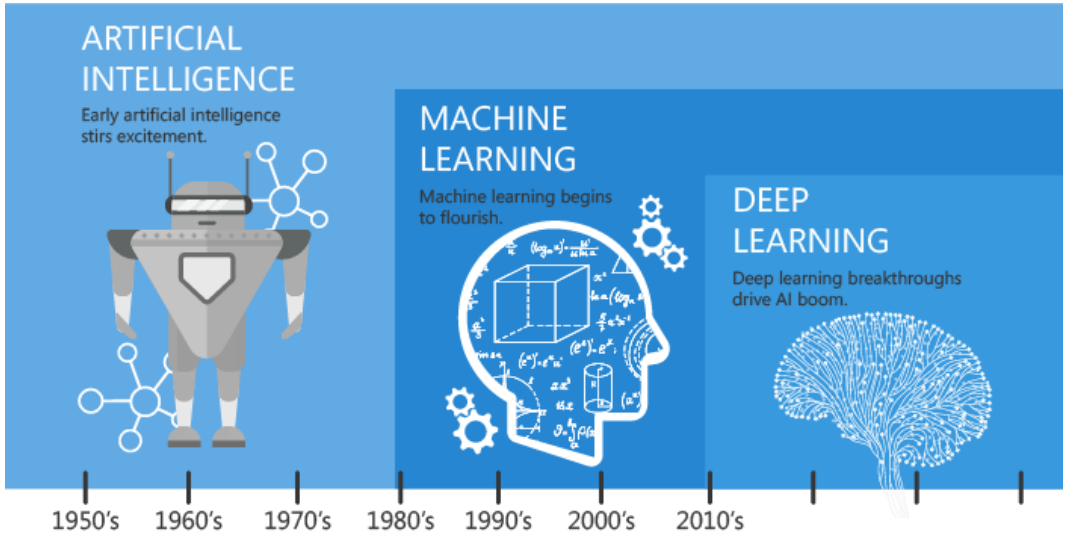

- Risks:
  - Bad data -> bad model
  - Unequal access to computation -> inquality (Timnit Gebru)
  - "black-box" (or maybe not… Increasingly debatable!)

Machine learning – how does it work?

- ***How do machines actually, learn?***

- Consider you are asked to find the best linear model that fits a set of data
  - And you skipped the linear algebra class

- But you know that: Y = mX + C
  - AKA, representation

- Parameters:
  - m - slope
  - C - bias

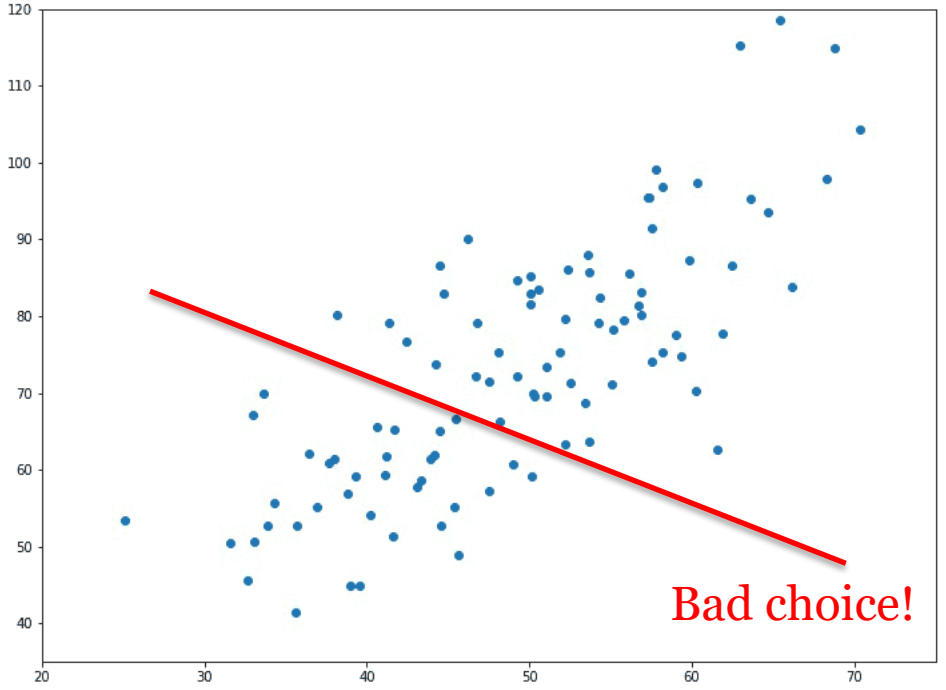

Machine learning – how does it work?

- ***How do machines actually, learn?***

- Consider you are asked to find the best linear model that fits a set of data
  - And you skipped the linear algebra class

- But you know that: Y = mX + C
  - AKA, representation

- Parameters:
  - m - slope
  - C - bias

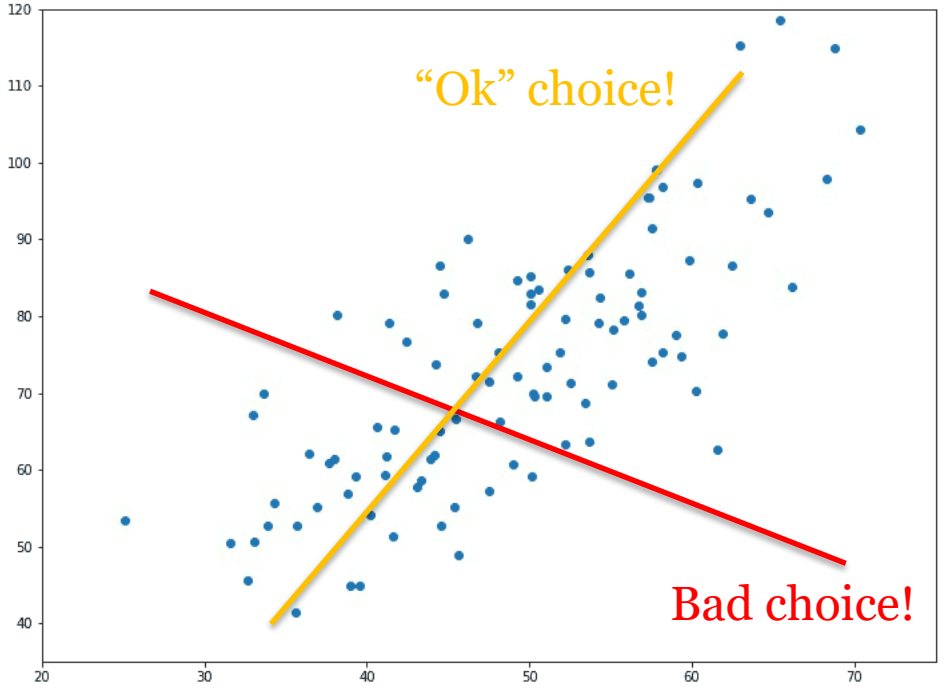

Machine learning – how does it work?

- ***How do machines actually, learn?***

- Consider you are asked to find the best linear model that fits a set of data
  - And you skipped the linear algebra class

- But you know that: Y = mX + C
  - AKA, representation

- Parameters:
  - m - slope
  - C - bias

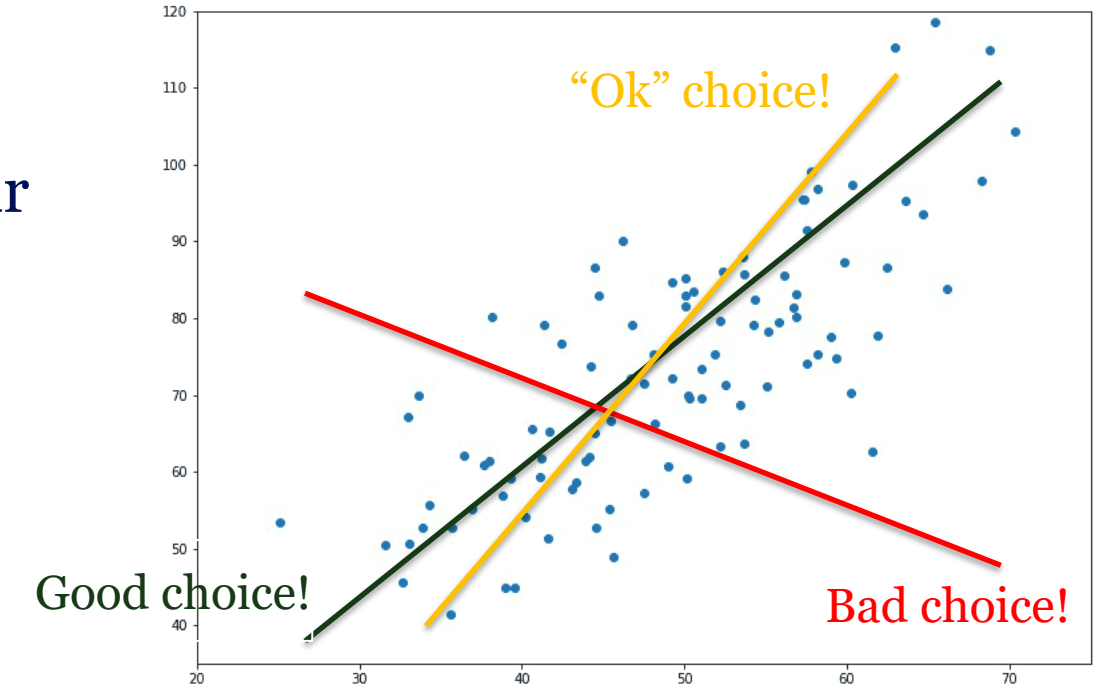

Evaluation:

$|| Y - Y_{red} || > || Y - Y_{yellow} || > || Y - Y_{green} ||$

Machine learning – how does it work?

3blue1brown: "But what is a neural network?" - Youtube playlist

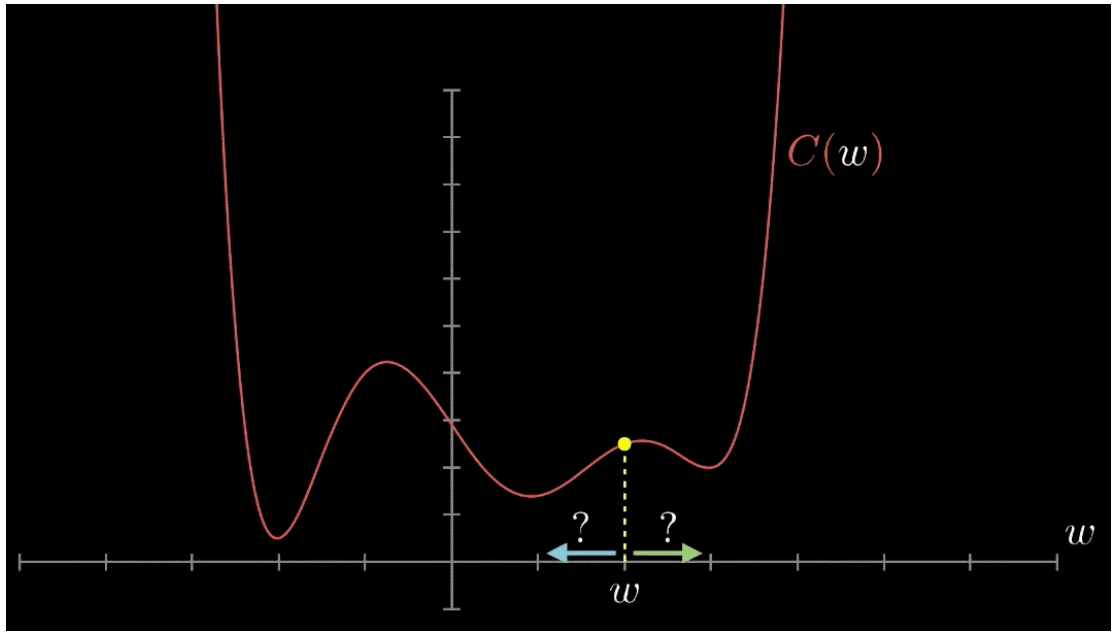

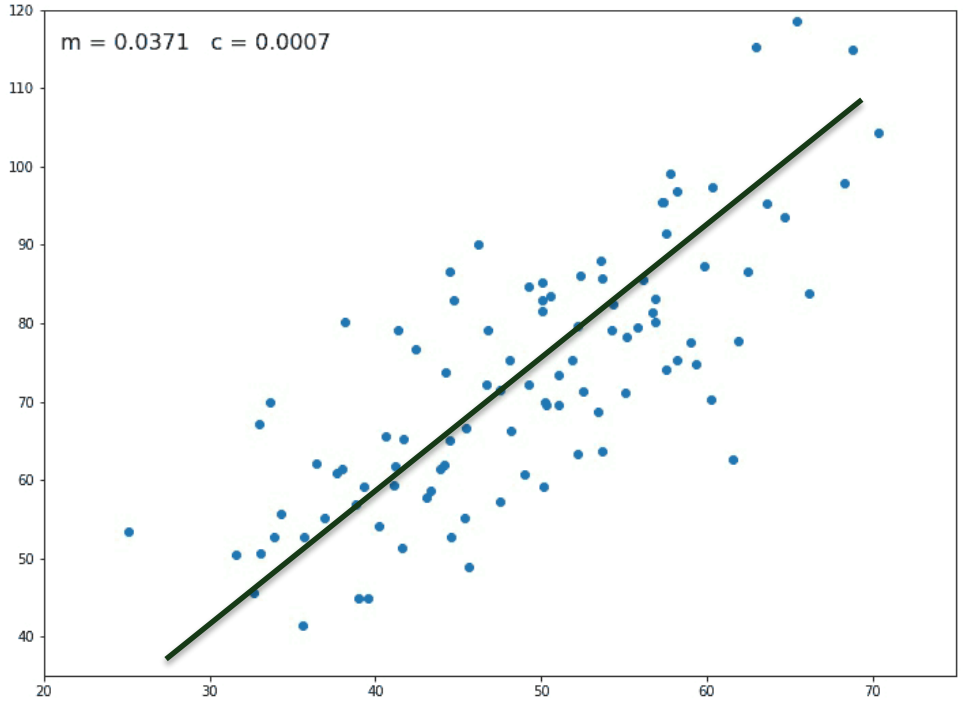

- If we would just try different random sets of parameters we might never find the best solution
- So, in general: start random but on each iteration, focus on the best results (optimization)

Machine learning – how does it work?

- At it's core:
  - Mapping functions: F(x) = y
  - The computer learns functions F(x) that transforms input X to output Y

$$F(\text{input}) = \text{Output}$$

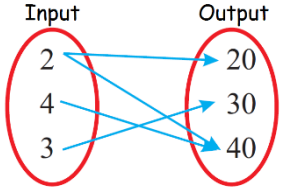

1. Algorithm/Representation
   - How to represent the relationship between the data

$$Y = mX + C$$

2. Evaluation
   - How to find if the model actually is working

$$||\ Y - Y_{red}\ || > ||\ Y - Y_{yellow}\ || > ||\ Y - Y_{green}\ ||$$

3. Optimization
   - How to search the possible solutions to find the best model

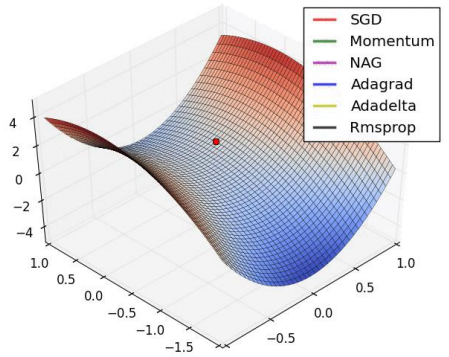

Any quick questions?

Universiteit Leiden
The Netherlands

Machine learning - Algorithms

- Random forest:
  - Decision trees (exactly what it means!)
  - Select rules that maximize similarity (e.g. Gini coefficient)
  - Fits Ordinary least squares regressions on the data subset

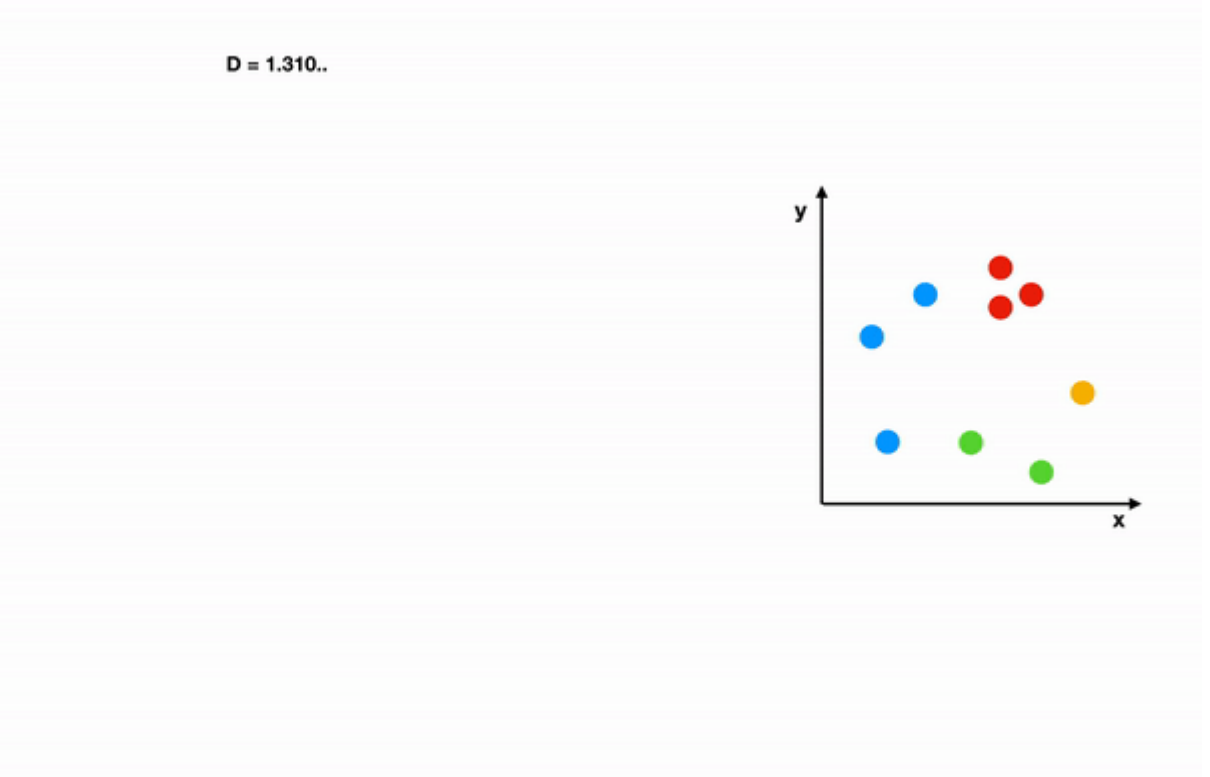

- Prediction:
  - Weighted sum of the weak models

- Typical parameters:
  - Maximum number of trees
  - Tree depth
  - Mininum number of points per “leaf” node, etc

Machine learning - Algorithms

- Support vector machines:
  - Separates classes by finding the best plane that maximizes the separabiity between classes
  - The plane is defined in the vector format (hence, support vector!)

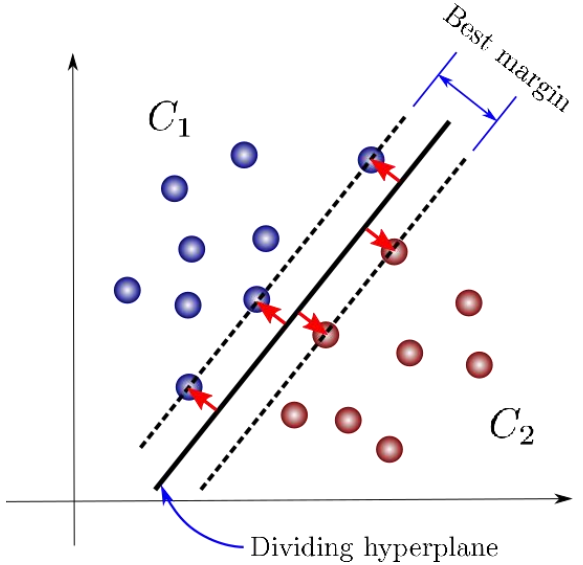

- For non linear behaviours - > kernel tricks
  - Extra feature spaces are created where linear hyperplanes can be created

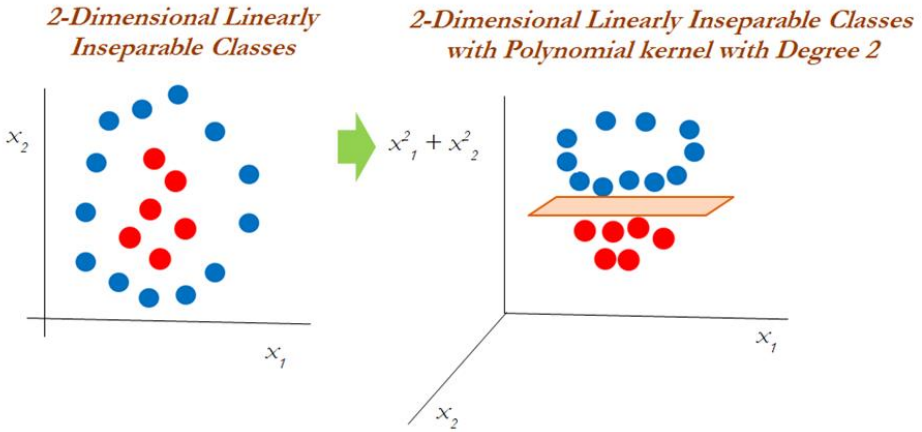

- Typical parameters:
  - C (regularization parameter); specific kernel parameters

$$K(X_1, X_2) = exp(-\frac{||X_1 - X_2||^2}{2\sigma^2})$$

Machine learning - Algorithms

- Artificial neural networks:
  - Initially developed to mimic human brain
  - Each neuron is a “linear model”
  - Activation function transforms linear models to non-linear functions

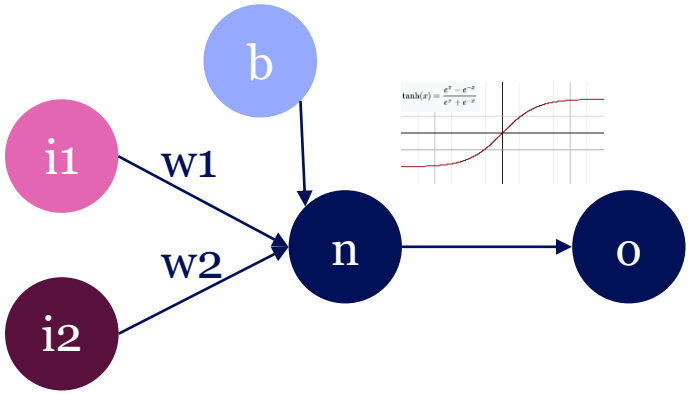

- Minimizes the error
  - On each iteration, small changes to the various w and b available

$$w1*i1 + w2*i2 + b = n$$

- Typical parameters:
  - How many neurons & layers?
  - What type of layers?
  - What activation functions?

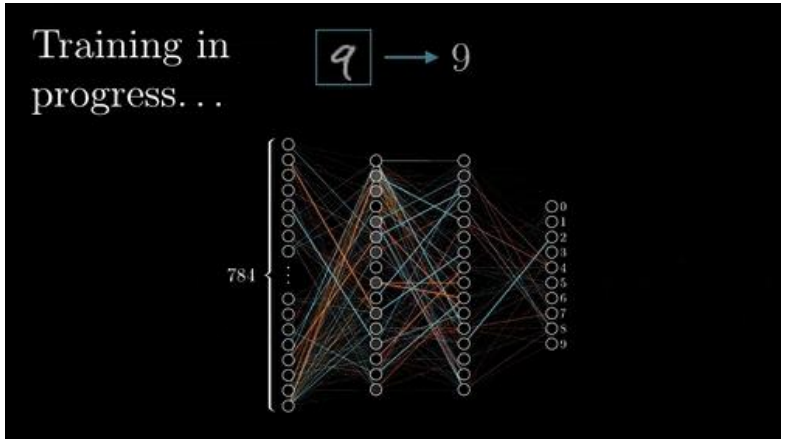

Machine learning – hyperparameters

- How many neurons and layers should a NN have?

- How to search efficiently for the model parameters?
  - We can “try” randomly and hope we find a good combination
  - We can be systematic
  - We can implement yet another optimization procedure..

- More advanced: Bayesian optimization

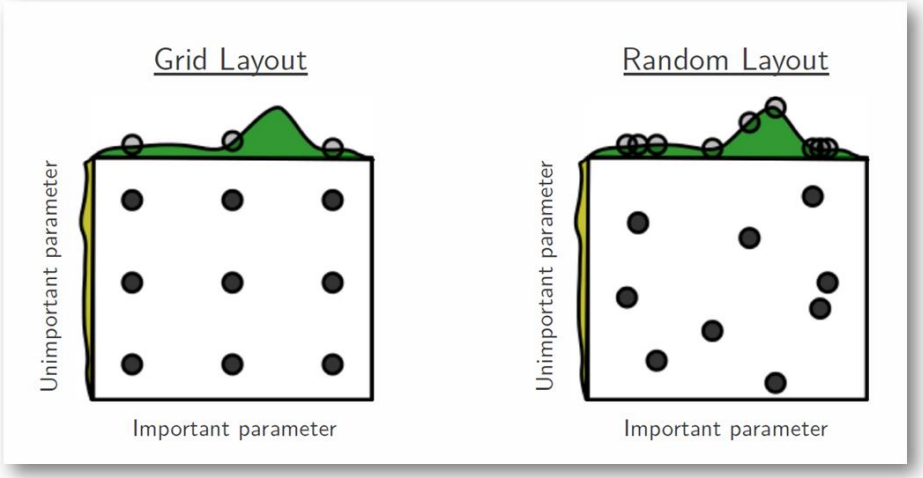

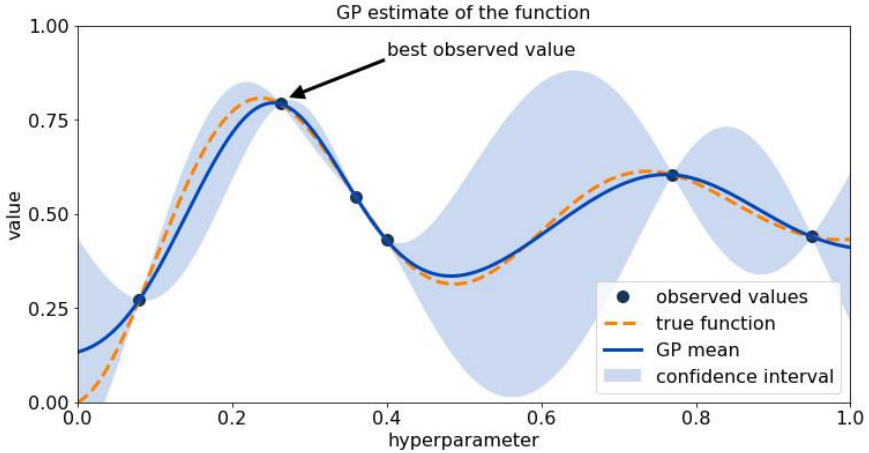

Machine learning – in species distributions

Ecological niche - EN

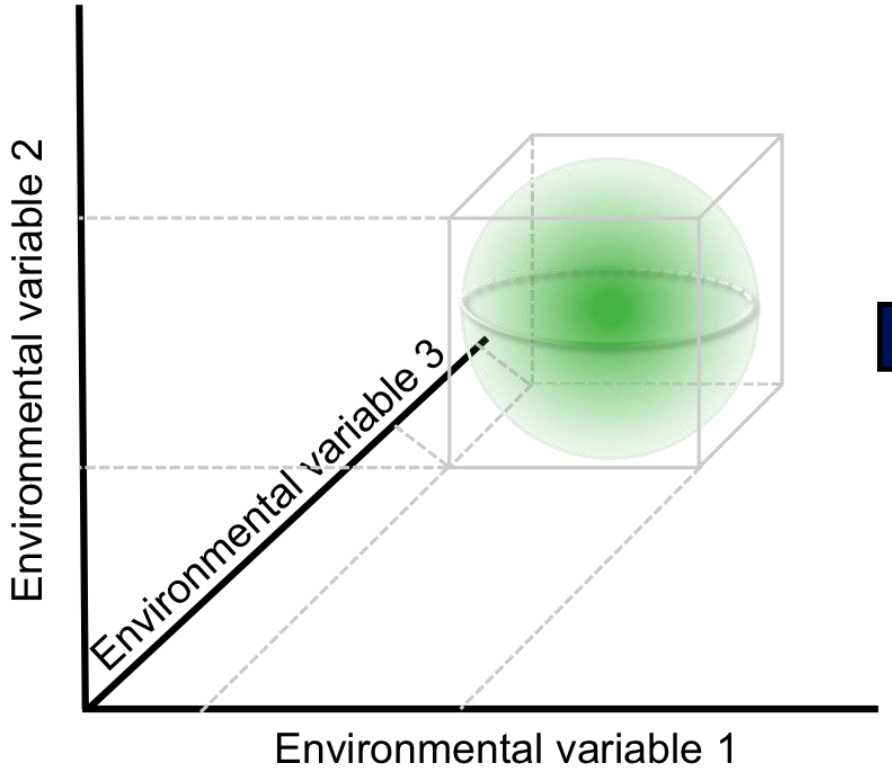

$f(E1, E2, E3)$

ML

Geographic space - EN projected by f(A,B,M)

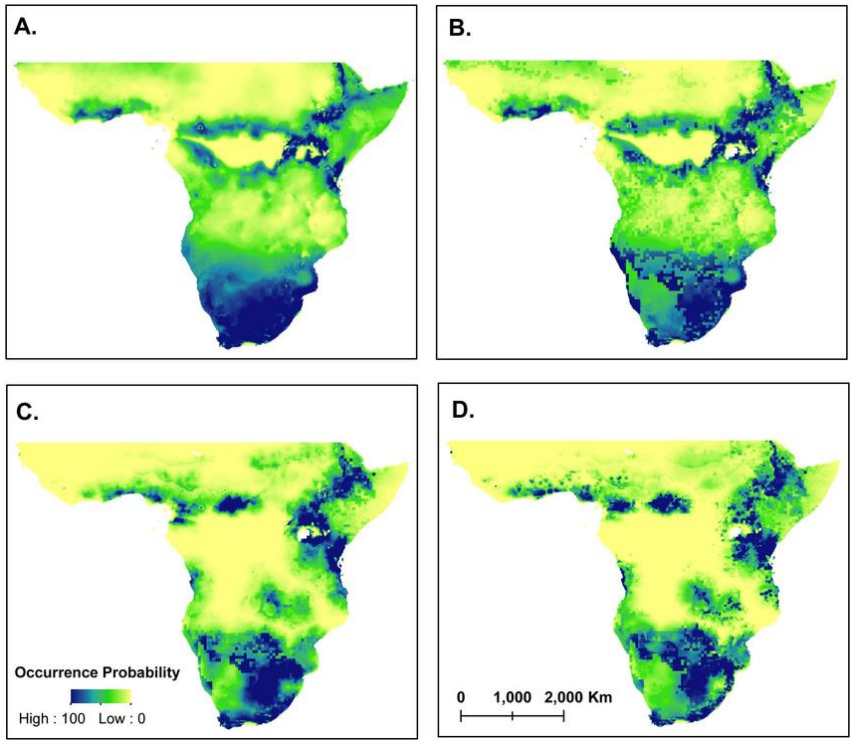

From: https://doi.org/10.4404/hystrix-27.1-11678

If the EN was a simple n-dimentional surface… It would be an easy problem….

Machine learning – In remote sensing

RS classification

Pixel-based

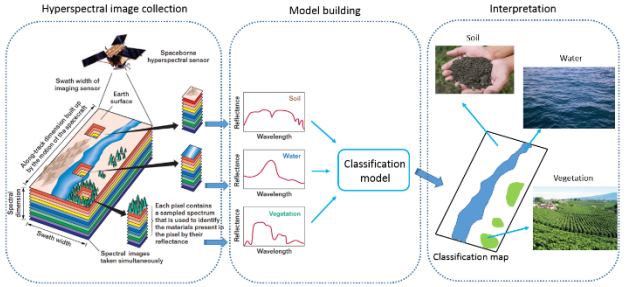

Object-based

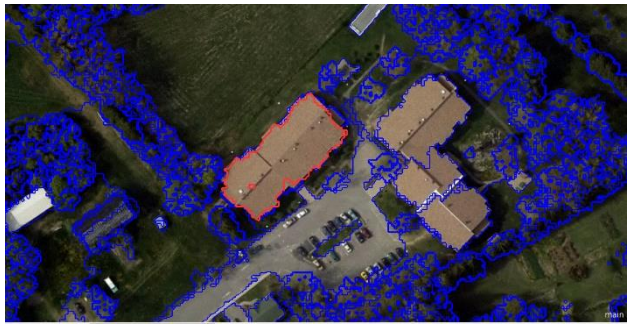

Computer vision

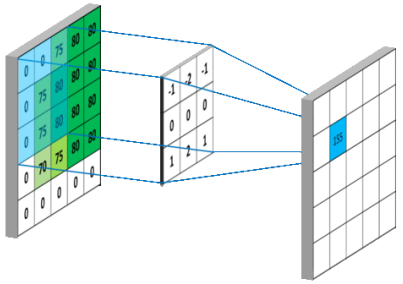

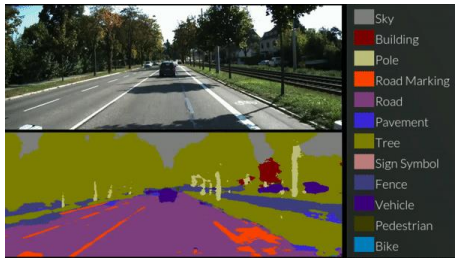

Any questions?

Universiteit Leiden
The Netherlands

Machine learning – where to start?

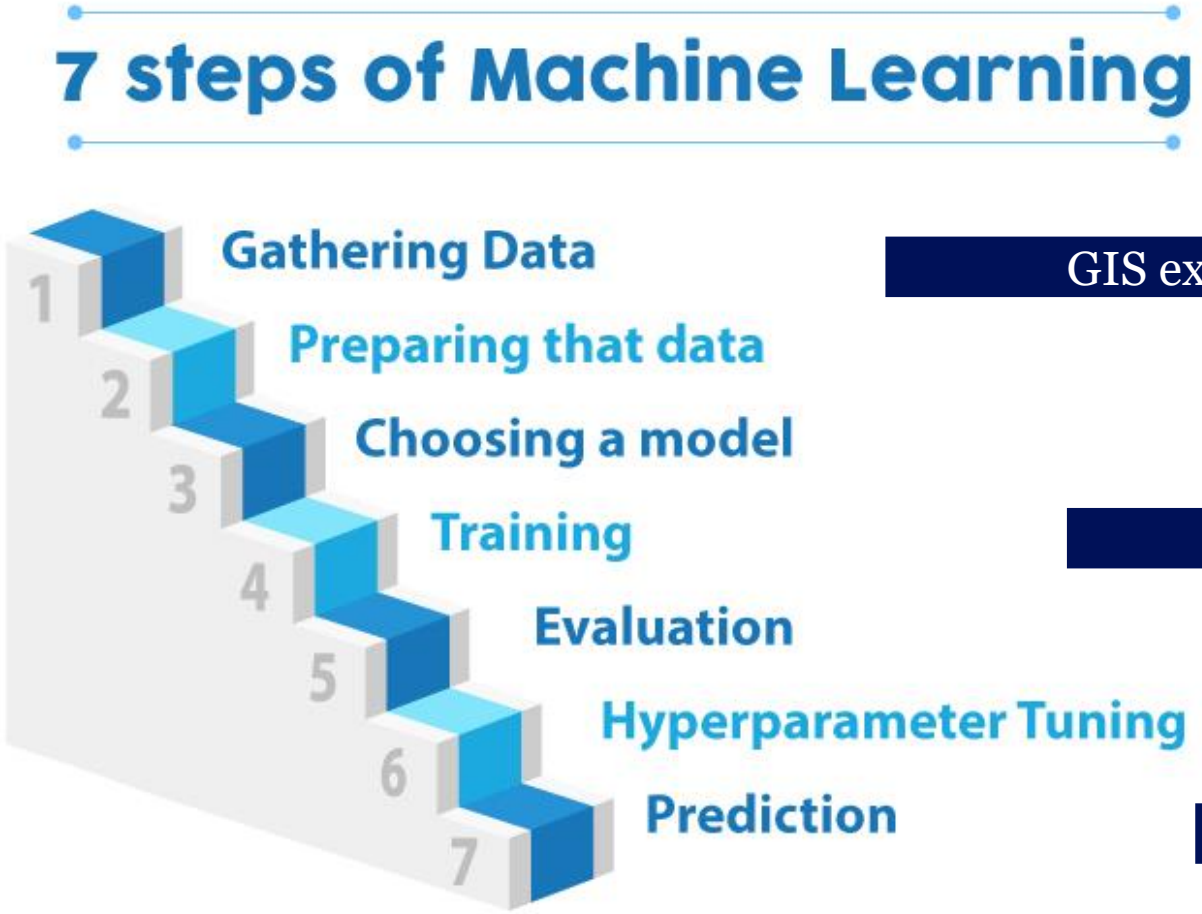

GIS expert

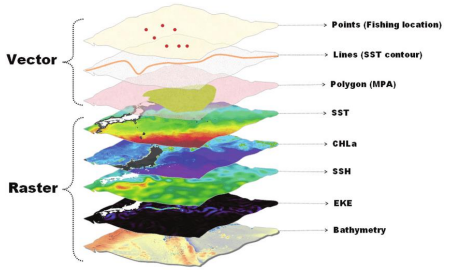

ML expert

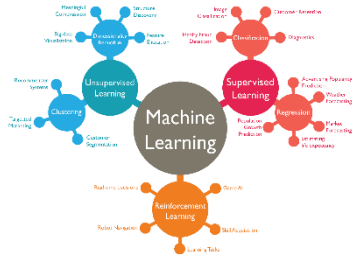

Client

The problem is:

- What is the best algorithm for that specific task?
- What are the best parameters?
- What pre-processing is required?
- How can “expertise” help?

All these need to be addressed.. If i want to scale up a business

AutoML

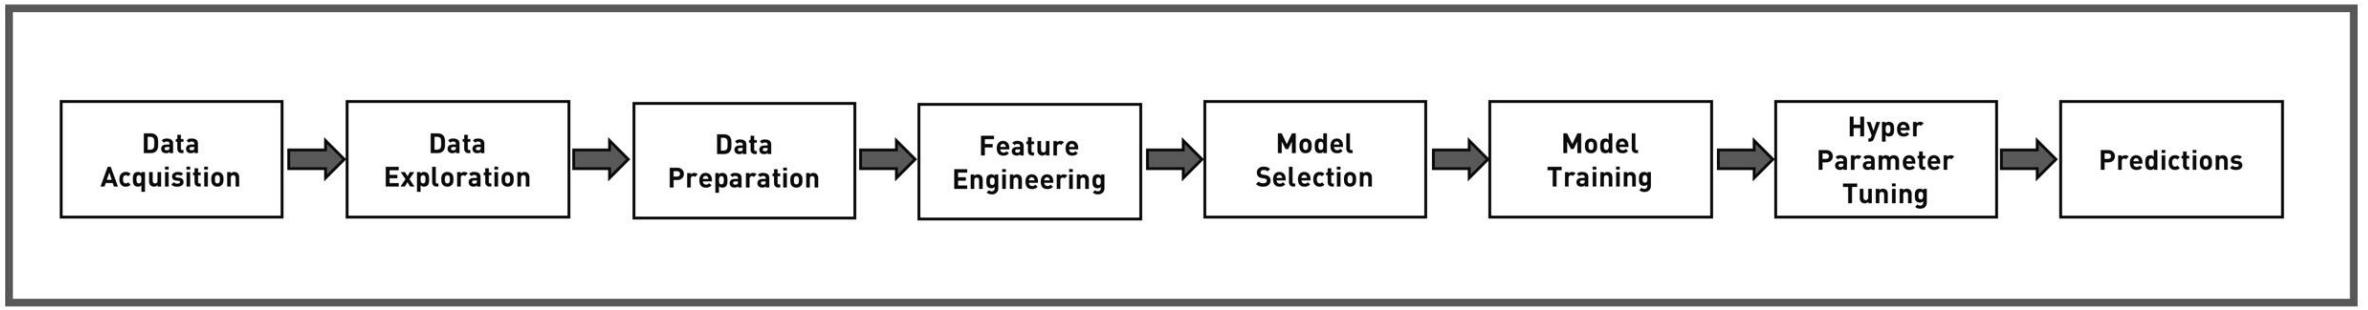

Traditional Machine Learning Workflow

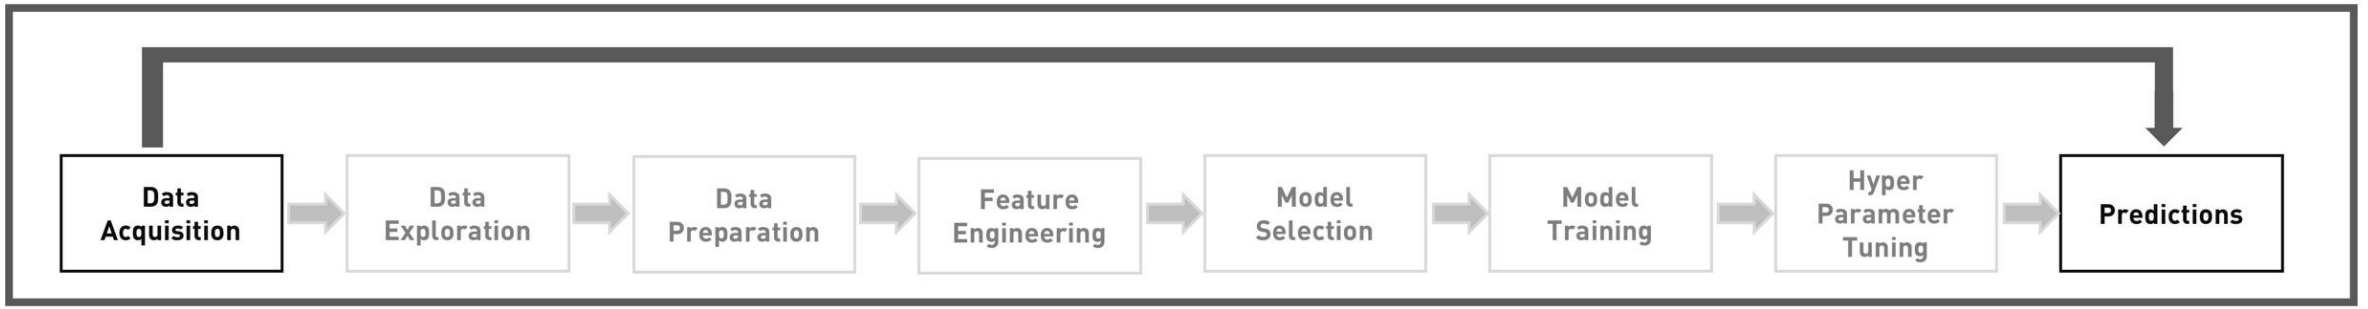

AutoML Workflow

AutoML

- Automation
  - Lower the entry level for machine learning
  - Optimal for deploying as a service
  - Minimum need to have an expert look at the clients data

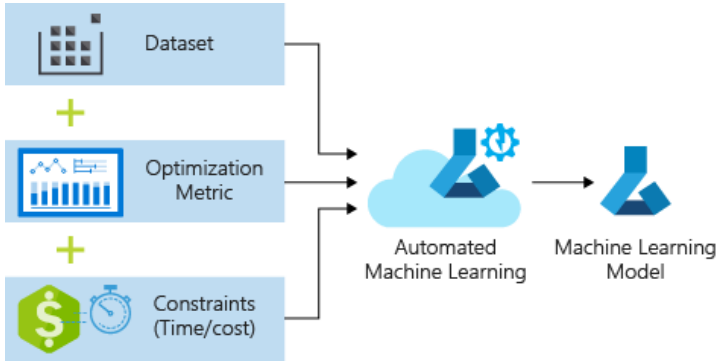

- CASH – Combined Algorithm Selection and Hyperparameter optimization

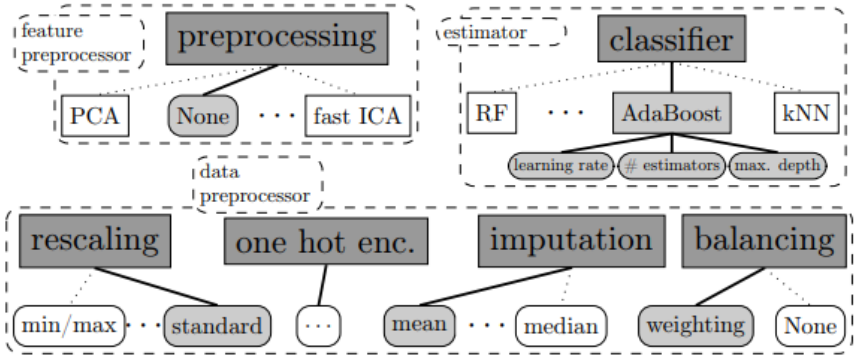

- Many options in the "market":
  - H2O (R also); Auto-Weka; TPOP; AutoGluon
  - Today: AutoSklearn

AutoML Framework	Open	Raw	Neural Network	CASH Strategy	Model Ensembling
Auto-WEKA	√	×	sigmoid MLP	Bayesian Optimization	bag, boost, stack, vote
auto-sklearn	√	×	none	BayesOpt + Meta-Learn	ensemble selection
TPOT	√	×	none	genetic programming	stacking
H2O	√	√	MLP + Adadelta	random search	stacking + bagging
GCP-Tables	×	√	AdaNet (??)	AdaNet (??)	boosting (??)
AutoGluon	√	√	embed categorical + skip-connection	fixed defaults (set adaptively)	multi-layer stacking + repeated bagging

AutoGluon-Tabular: Robbust and Accurate AutoML for Structured data, 2020

AutoSklearn

- Uses the algorithms available in scikit-learn
  - The “standard” python library for Machine learning
  - Many different algorithms available

- CASH:
  - Bayesian optimization
  - Meta-learning

WE DID IT AGAIN: WORLD CHAMPIONS IN AUTOML

- Meta-learning:
  - Pre-tested pipelines first (emulates expertise)
  - Based on pre-trained pipelines on open data

https://www.automl.org/

AutoSklearn

- Pre-processing of data and features
- Algorithm selection

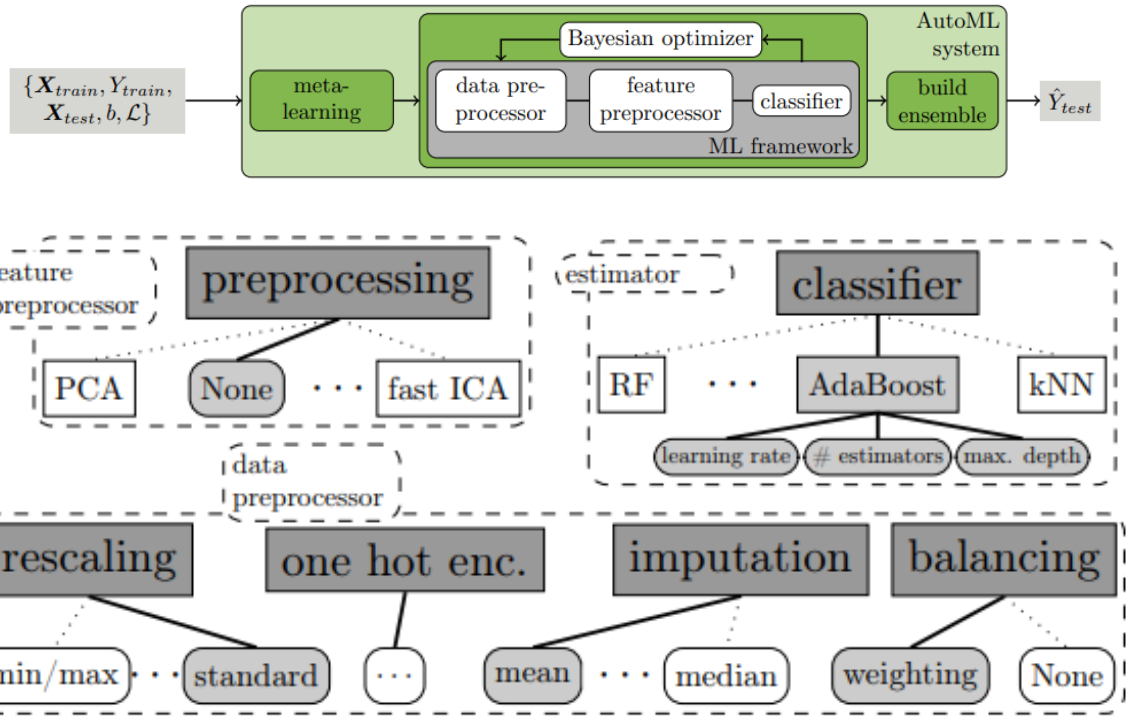

[Efficient and Robust Automated ML](), Feurer et al, 2015

Bayesian optimization

- Uses a probabilistic approach to find the "options" that promises most improvement

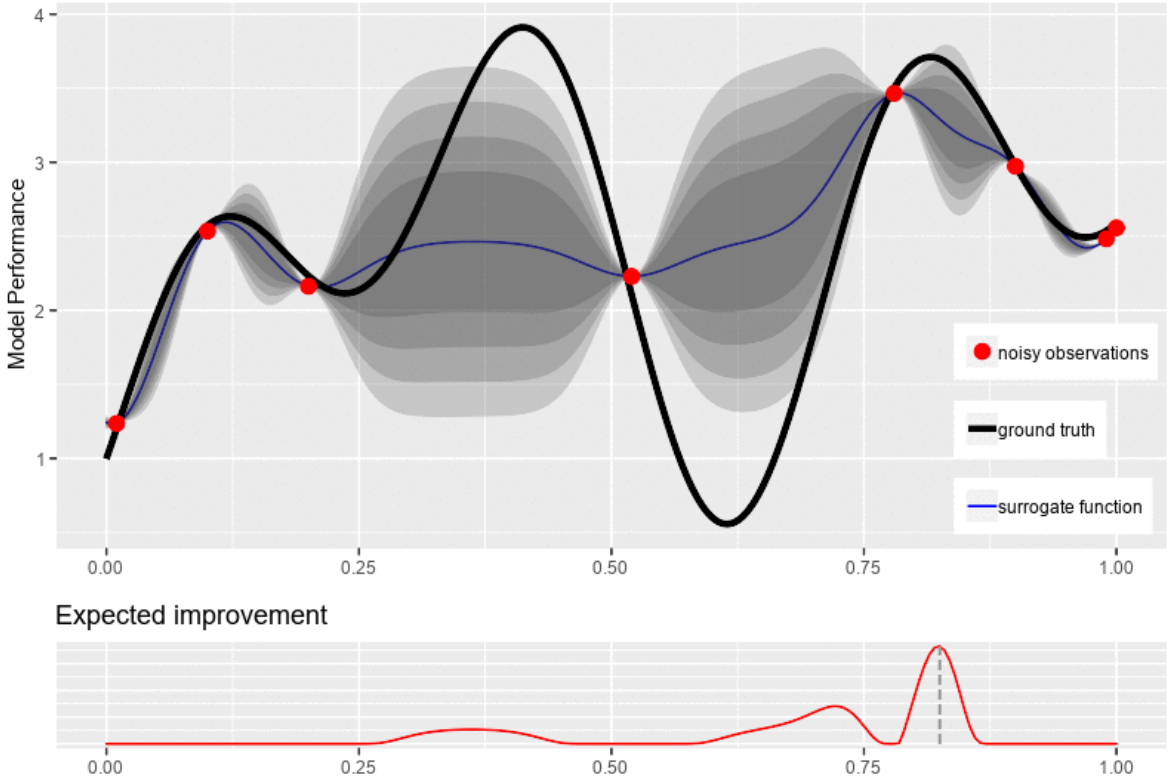

Classification example:

- Applying to RS data:
  - Spectral classification
  - Adapted from Sydney Goldstein tutorial on R

- Landsat 8 image, Calgary
  - Clouds
  - Developed
  - Underdeveloped
  - Water

- Accuracy:
  - Overall: 0.9186
  - Kappa: 0.89

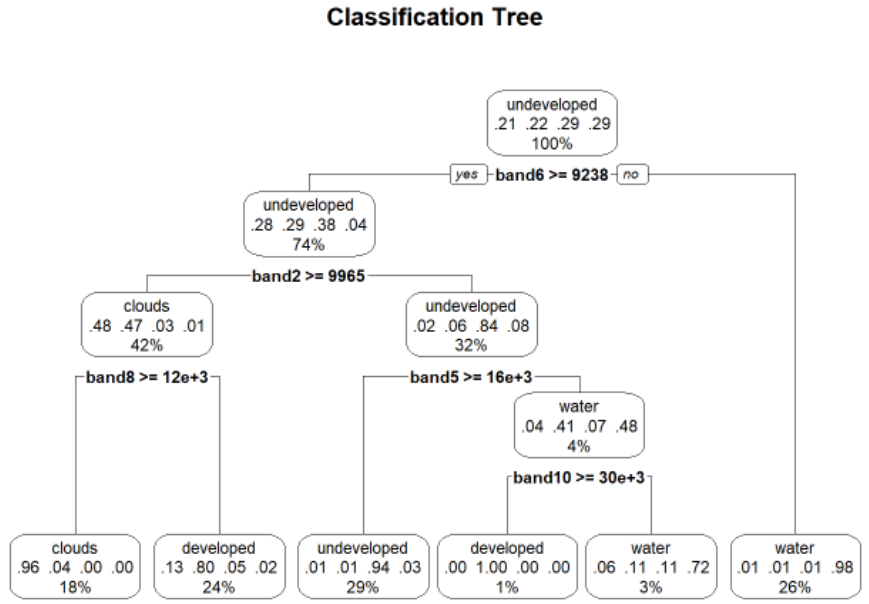

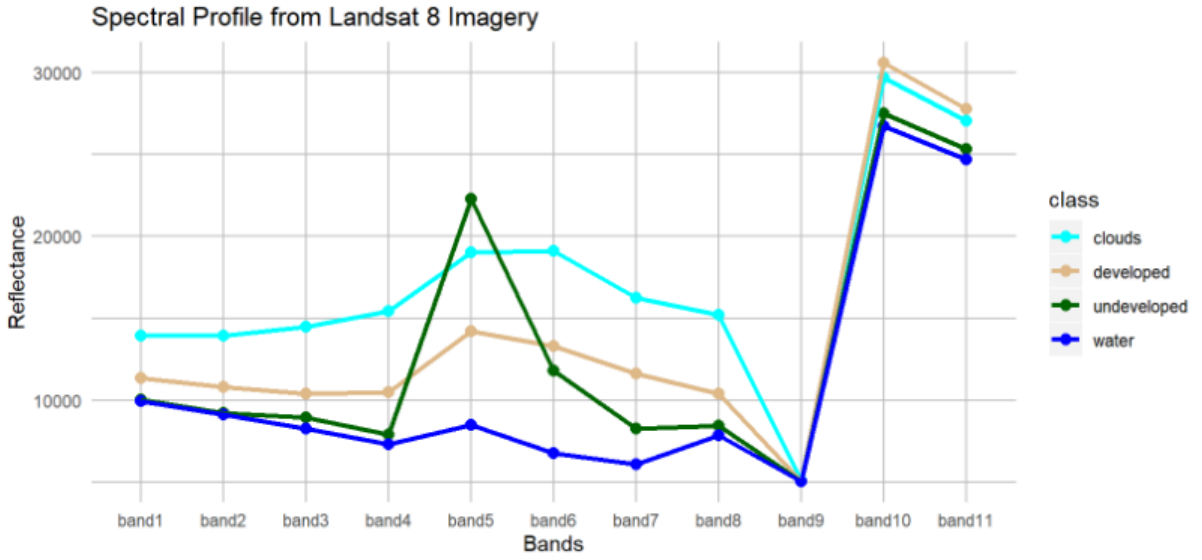

AutoSklearn on RS data

- Only bands 1,2,3,4,5,6,7
  - I was lazy to deal with the resolution of the rest..
- AutoSklearn configurations:
  - Use meta-learning
  - Max time 60s
  - 5-fold crossvalidation
- [Github: nunocesarsa/Examples/AutoML/AutoSklearn]
- Accuracy: ~ 0.96

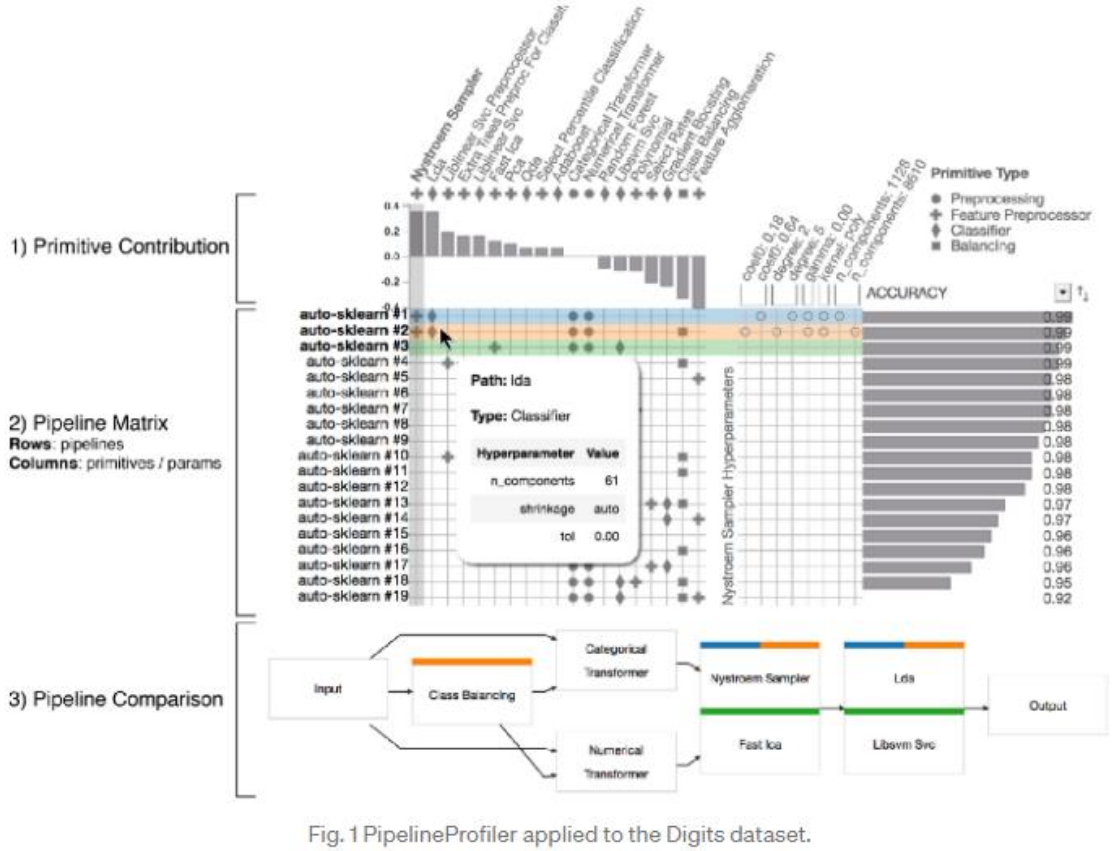

Fig. 1 PipelineProfiler applied to the Digits dataset.

https://towardsdatascience.com/exploring-auto-sklearn-models-with-pipelineprofiler-5b2c54136044

AutoGluon

- Developed by AWS to be deployed as service
  - Work for very different tasks (images, tables, text, NAS)
- Different: not a CASH approach
  - No hyperparameter tunning (possible, but not recommended)
- "Multi-layer stacking"
  - On each layer, the predictions of the models on the previous layers are added as features to the next layer
- Weak learners combined > strong learners

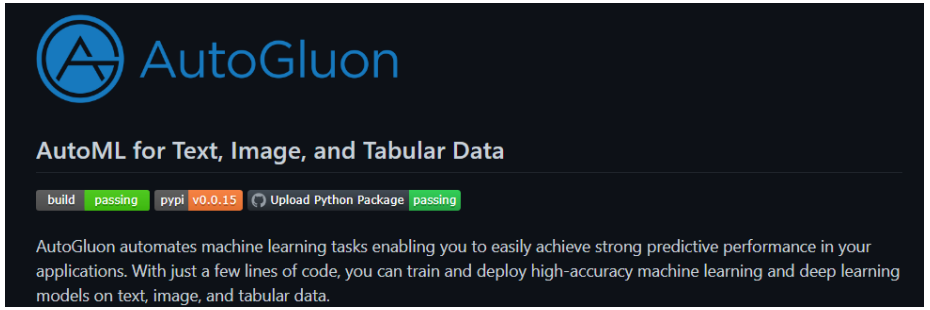

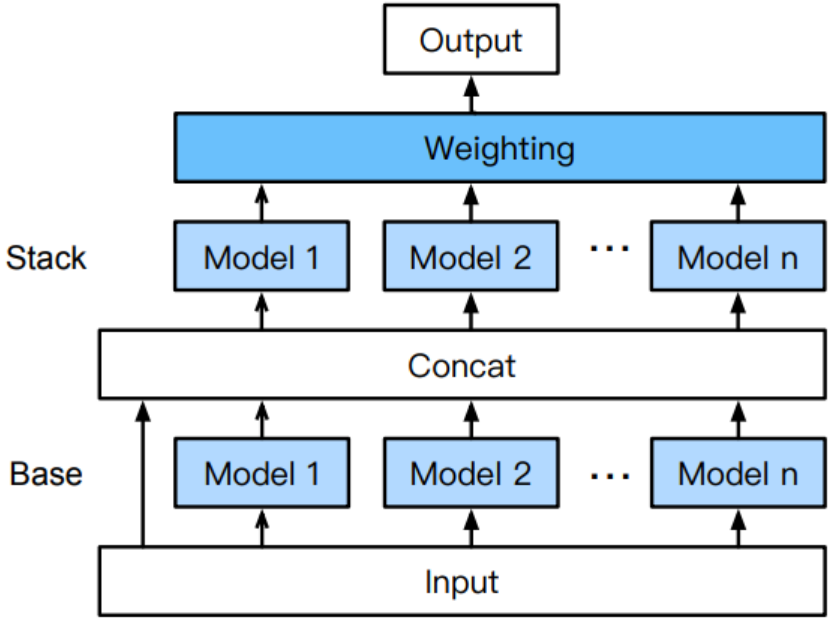

AutoGluon

- There are a number of “pre-sets” that are changeable:
- In our case, we opted for a very simple preset:
  - ‘Hyperameter tuning: light -> meaning very small tunning
- Sequential model training
- Accuracy:
  - Overall: 0.95
  - Kappa: 0.93

```
Train Data Columns: 7
Preprocessing data ...
AutoGluon infers your prediction problem is: 'multiclass' (because dtype of label-column == category).
        4 unique label values:  ['undeveloped', 'water', 'developed', 'clouds']
        If 'multiclass' is not the correct problem_type, please manually specify the problem_type argume
Train Data Class Count: 4
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
        Available Memory:                    12473.82 MB
        Train Data (Original)  Memory Usage: 0.03 MB (0.0% of available memory)
        Inferring data type of each feature based on column values. Set feature_metadata_in to manually
        Stage 1 Generators:
```

```
        To change this, specify the eval_metric argument of fit()
AutoGluon will early stop models using evaluation metric: 'accuracy'
Fitting model: NeuralNetClassifier_STACKER_l0 ... Training model for up to 179.91s of the
        0.9327   = Validation accuracy score
        9.51s    = Training runtime
        0.06s    = Validation runtime
Fitting model: RandomForestClassifierGini_STACKER_l0 ... Training model for up to 170.3s o
        0.9245   = Validation accuracy score
        3.46s    = Training runtime
        0.51s    = Validation runtime
Fitting model: RandomForestClassifierEntr_STACKER_l0 ... Training model for up to 166.25s
        0.9224   = Validation accuracy score
```

Github: https://github.com/nunocesarsa/Examples

In summary:

Automated machine learning:

1. Improves accuracy & reduces the “human-in-the-loop”
   - Critical criteria for any cloud service

2. Can provide more insights than just single model approaches
   - Not show.. Autosklearn pipeliner is very cool

3. Can/should be fine-tuned for your tasks
   - (e.g., meta-learning from RS classifications)
   - E.g., in the RS field would be benefit from a Vegetation indice feature generator?

4. Deployable!
   - And with pre-set time.. And quicker than any “hands-on” approach

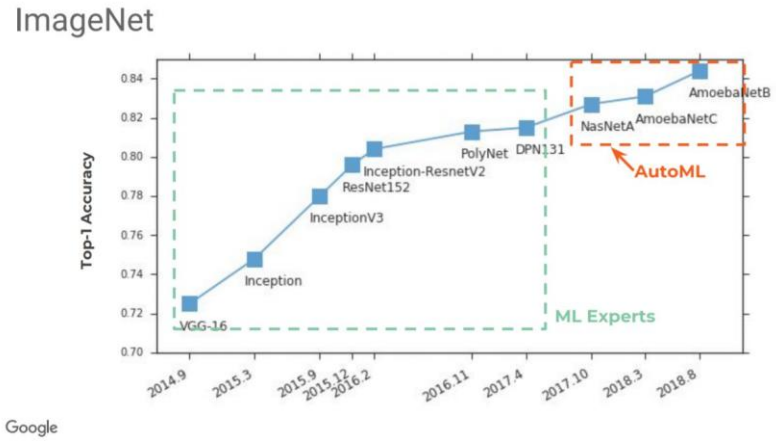

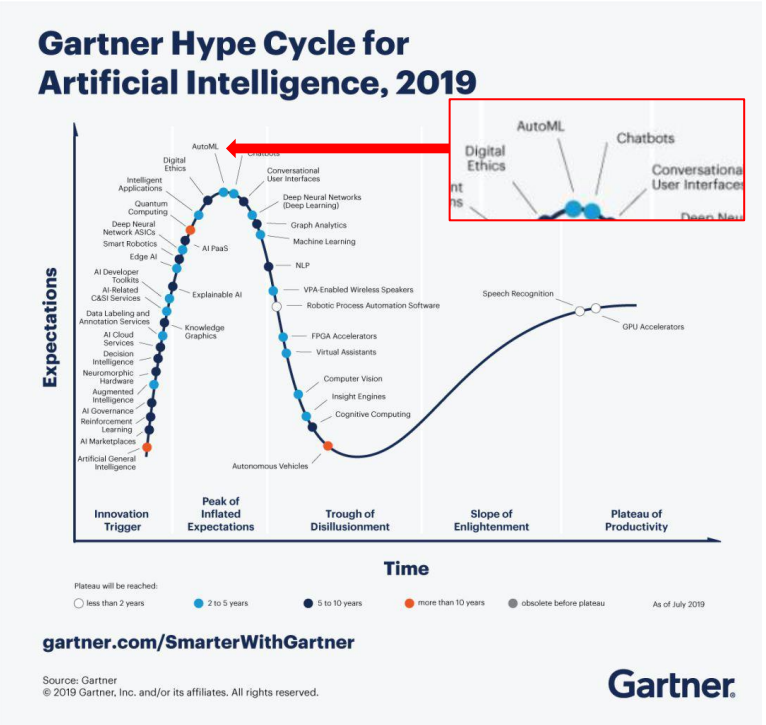

Any questions?

Universiteit Leiden
The Netherlands

We can go for the colab after/during